

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
from matplotlib.pyplot import rcParams
rcParams['figure.figsize'] = 10, 6
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: dataset=pd.read_excel("C:\\Users\\Admin\\Downloads\\assignment 10\\Airlines+Data.xlsx")
dataset.head()
```

```
Out[2]:
```

	Month	Passengers
0	1995-01-01	112
1	1995-02-01	118
2	1995-03-01	132
3	1995-04-01	129
4	1995-05-01	121

```
In [3]: dataset['Month'] = pd.to_datetime(dataset['Month'], infer_datetime_format=True)
indexedDataset = dataset.set_index(['Month'])
```

```
In [4]: from datetime import datetime
indexedDataset['1995-03']
indexedDataset['1995-03':'1995-06']
indexedDataset['1995']
```


```
Out[4]:
```

	Passengers
Month	
1995-01-01	112
1995-02-01	118
1995-03-01	132
1995-04-01	129
1995-05-01	121
1995-06-01	135
1995-07-01	148
1995-08-01	148
1995-09-01	136
1995-10-01	119
1995-11-01	104

Passengers	
Month	
1995-12-01	118

```
In [5]: plt.xlabel("Date")
plt.ylabel("Number of air passengers")
plt.plot(indexedDataset)
```

Out[5]: [



```
In [6]: rolmean = indexedDataset.rolling(window=12).mean()
rolstd = indexedDataset.rolling(window=12).std()
print(rolmean, rolstd)
```

Passengers	
Month	
1995-01-01	NaN
1995-02-01	NaN
1995-03-01	NaN
1995-04-01	NaN
1995-05-01	NaN
...	...
2002-08-01	316.833333
2002-09-01	320.416667
2002-10-01	323.083333
2002-11-01	325.916667
2002-12-01	328.250000

[96 rows x 1 columns]		Passengers
Month		
1995-01-01	NaN	
1995-02-01	NaN	

```

1995-03-01      NaN
1995-04-01      NaN
1995-05-01      NaN
...
2002-08-01    54.530781
2002-09-01    55.586883
2002-10-01    53.899668
2002-11-01    49.692616
2002-12-01    47.861780


```

```
[96 rows x 1 columns]
```

```

In [7]: orig = plt.plot(indexedDataset, color='blue', label='Original')
mean = plt.plot(rolmean, color='red', label='Rolling Mean')
std = plt.plot(rolstd, color='black', label = 'Rolling Std')
plt.legend(loc='best')
plt.title('Rolling Mean & Standard Deviation')
plt.show(block=False)

```



```

In [8]: from statsmodels.tsa.stattools import adfuller

print ('Results of Dickey-Fuller Test:')
dfctest = adfuller(indexedDataset['Passengers'], autolag='AIC')
dfcoutput = pd.Series(dfctest[0:4], index=['Test Statistic', 'p-value', '#Lags Used', 'Number of Observations Used'])
for key,value in dfctest[4].items():
    dfcoutput['Critical Value (%s)'%key] = value
print(dfcoutput)

```

```


Results of Dickey-Fuller Test:
Test Statistic      1.340248
p-value             0.996825
#Lags Used          12.000000
Number of Observations Used  83.000000
Critical Value (1%) -3.511712

```

```
Critical Value (5%)          -2.897048  
Critical Value (10%)         -2.585713  
dtype: float64
```


```
In [9]: indexedDataset_logScale = np.log(indexedDataset)  
plt.plot(indexedDataset_logScale)
```

```
Out[9]: [<matplotlib.lines.Line2D at 0x203bcf002b0>]
```



```
In [10]: movingAverage = indexedDataset_logScale.rolling(window=12).mean()  
movingSTD = indexedDataset_logScale.rolling(window=12).std()  
plt.plot(indexedDataset_logScale)  
plt.plot(movingAverage, color='red')
```

```
Out[10]: [<matplotlib.lines.Line2D at 0x203bd51d190>]
```



```
In [11]: datasetLogScaleMinusMovingAverage = indexedDataset_logScale - movingAverage
datasetLogScaleMinusMovingAverage.head(12)
#Remove Nan Values
datasetLogScaleMinusMovingAverage.dropna(inplace=True)
datasetLogScaleMinusMovingAverage.head(10)
```

Out[11]: **Passengers**

Month	
1995-12-01	-0.065494
1996-01-01	-0.093449
1996-02-01	-0.007566
1996-03-01	0.099416
1996-04-01	0.052142
1996-05-01	-0.027529
1996-06-01	0.139881
1996-07-01	0.260184
1996-08-01	0.248635
1996-09-01	0.162937


```
In [12]: from statsmodels.tsa.stattools import adfuller
def test_stationarity(timeseries):

    #Determining rolling statistics
    movingAverage = timeseries.rolling(window=12).mean()
    movingSTD = timeseries.rolling(window=12).std()
```

```
#Plot rolling statistics:
orig = plt.plot(timeseries, color='blue',label='Original')
mean = plt.plot(movingAverage, color='red', label='Rolling Mean')
std = plt.plot(movingSTD, color='black', label = 'Rolling Std')
plt.legend(loc='best')
plt.title('Rolling Mean & Standard Deviation')
plt.show(block=False)

#Perform Dickey-Fuller test:
print('Results of Dickey-Fuller Test:')
dfctest = adfuller(timeseries['Passengers'], autolag='AIC')
dfcoutput = pd.Series(dfctest[0:4], index=['Test Statistic','p-value','#Lags Used','N
for key,value in dfctest[4].items():
    dfcoutput['Critical Value (%)'%key] = value
print(dfcoutput)
```

In [13]: test_stationarity(datasetLogScaleMinusMovingAverage)




```
Results of Dickey-Fuller Test:
Test Statistic      -1.910930
p-value             0.326937
#Lags Used          12.000000
Number of Observations Used  72.000000
Critical Value (1%)   -3.524624
Critical Value (5%)   -2.902607
Critical Value (10%)  -2.588679
dtype: float64
```


In [14]:

```
exponentialDecayWeightedAverage = indexedDataset_logScale.ewm(halflife=12, min_periods=
plt.plot(indexedDataset_logScale)
plt.plot(exponentialDecayWeightedAverage, color='red')
```

Out[14]: [<matplotlib.lines.Line2D at 0x203bd498f70>]



```
In [15]: datasetLogScaleMinusMovingExponentialDecayAverage = indexedDataset_logScale - exponenti
test_stationarity(datasetLogScaleMinusMovingExponentialDecayAverage)
```




Results of Dickey-Fuller Test:


Test Statistic	-2.835036
p-value	0.053441
#Lags Used	12.000000
Number of Observations Used	83.000000
Critical Value (1%)	-3.511712
Critical Value (5%)	-2.897048
Critical Value (10%)	-2.585713
dtype:	float64

```
In [16]: datasetLogDiffShifting = indexedDataset_logScale - indexedDataset_logScale.shift()  
plt.plot(datasetLogDiffShifting)
```

```
Out[16]: [<matplotlib.lines.Line2D at 0x203bd5714f0>]
```



```
In [17]: datasetLogDiffShifting.dropna(inplace=True)  
test_stationarity(datasetLogDiffShifting)
```



Results of Dickey-Fuller Test:

Test Statistic	-2.670823
p-value	0.079225
#Lags Used	12.000000
Number of Observations Used	82.000000

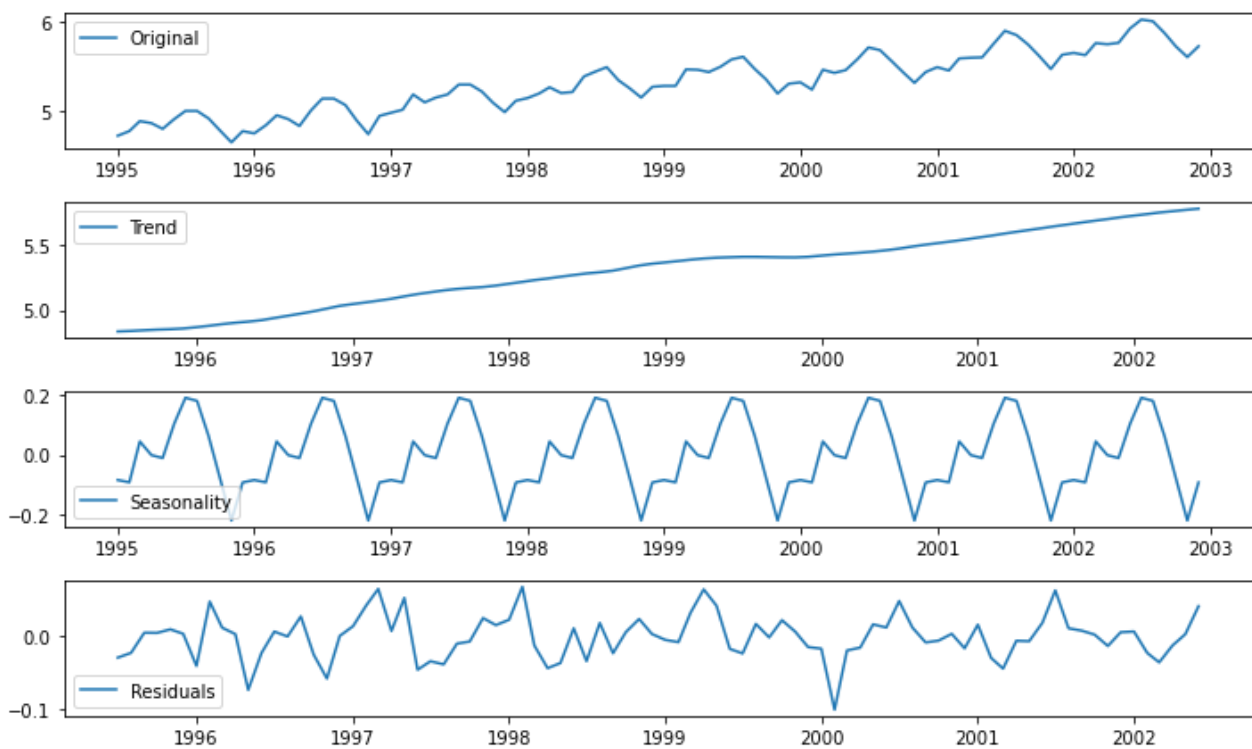
Critical Value (1%) -3.512738
 Critical Value (5%) -2.897490
 Critical Value (10%) -2.585949
 dtype: float64

In [18]:

```
from statsmodels.tsa.seasonal import seasonal_decompose
decomposition = seasonal_decompose(indexedDataset_logScale)

trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid

plt.subplot(411)
plt.plot(indexedDataset_logScale, label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend, label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal, label='Seasonality')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual, label='Residuals')
plt.legend(loc='best')
plt.tight_layout()
```



In [19]:


```
from statsmodels.tsa.stattools import acf, pacf

lag_acf = acf(datasetLogDiffShifting, nlags=20)
lag_pacf = pacf(datasetLogDiffShifting, nlags=20, method='ols')

#Plot ACF:
plt.subplot(121)
```

```
plt.plot(lag_acf)
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(datasetLogDiffShifting)),linestyle='--',color='gray')
plt.axhline(y=1.96/np.sqrt(len(datasetLogDiffShifting)),linestyle='--',color='gray')
plt.title('Autocorrelation Function')

#Plot PACF:
plt.subplot(122)
plt.plot(lag_pacf)
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(datasetLogDiffShifting)),linestyle='--',color='gray')
plt.axhline(y=1.96/np.sqrt(len(datasetLogDiffShifting)),linestyle='--',color='gray')
plt.title('Partial Autocorrelation Function')
plt.tight_layout()
```

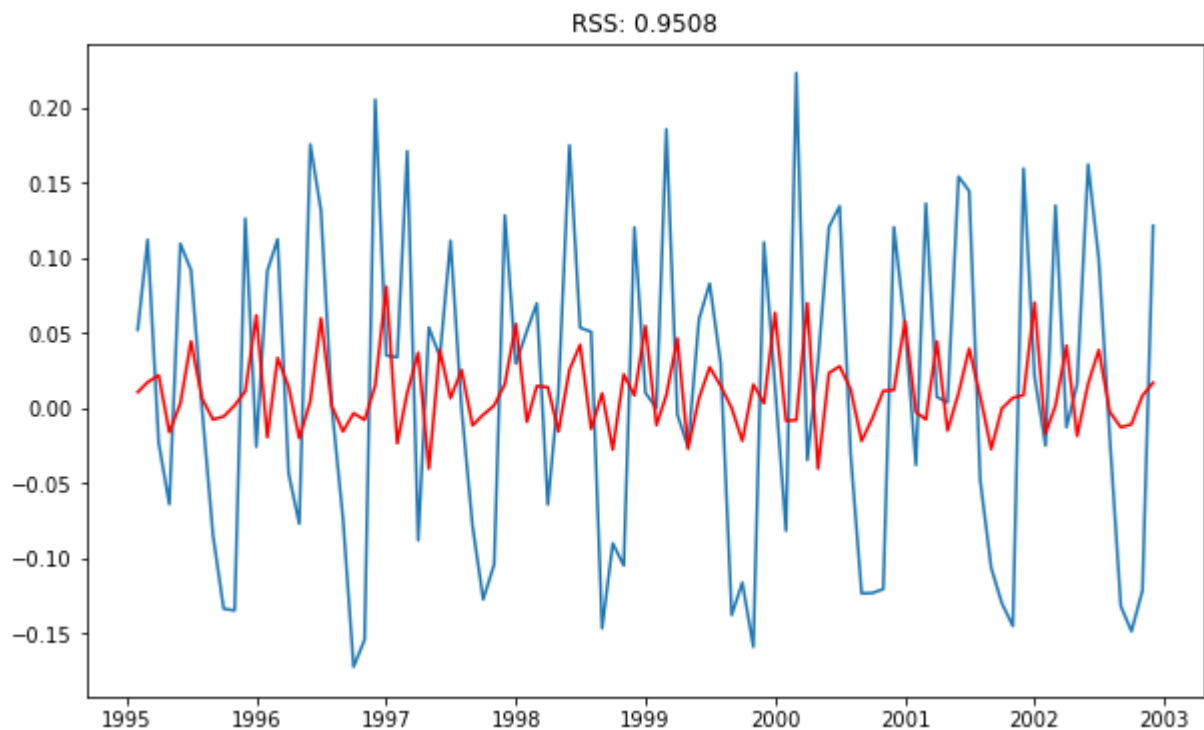


In [20]:

```
from statsmodels.tsa.arima_model import ARIMA

#AR MODEL
model = ARIMA(indexedDataset_logScale, order=(2, 1, 0))
results_AR = model.fit(disp=-1)
plt.plot(datasetLogDiffShifting)
plt.plot(results_AR.fittedvalues, color='red')
plt.title('RSS: %.4f%% sum((results_AR.fittedvalues-datasetLogDiffShifting["Passengers"]
print('Plotting AR model')
```

C:\Users\Admin\anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
 warnings.warn('No frequency information was'
C:\Users\Admin\anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
 warnings.warn('No frequency information was'
Plotting AR model



In [21]:

```

model = ARIMA(indexedDataset_logScale, order=(0, 1, 2))
results_MA = model.fit(dis=-1)
plt.plot(datasetLogDiffShifting)
plt.plot(results_MA.fittedvalues, color='red')
plt.title('RSS: %.4f'% sum((results_MA.fittedvalues-datasetLogDiffShifting["Passengers"]
print('Plotting AR model')

```


C:\Users\Admin\anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

warnings.warn('No frequency information was')

C:\Users\Admin\anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

warnings.warn('No frequency information was')

Plotting AR model



```
In [22]: model = ARIMA(indexedDataset_logScale, order=(2, 1, 2))
results_ARIMA = model.fit(dispatch=-1)
plt.plot(datasetLogDiffShifting)
plt.plot(results_ARIMA.fittedvalues, color='red')
plt.title('RSS: %.4f'% sum((results_ARIMA.fittedvalues-datasetLogDiffShifting["Passenge
```


C:\Users\Admin\anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

warnings.warn('No frequency information was'

C:\Users\Admin\anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

warnings.warn('No frequency information was'

```
Out[22]: Text(0.5, 1.0, 'RSS: 0.6931')
```



```
In [23]: predictions_ARIMA_diff = pd.Series(results_ARIMA.fittedvalues, copy=True)
print (predictions_ARIMA_diff.head())
```

```
Month
1995-02-01    0.011261
1995-03-01    0.016602
1995-04-01    0.021663
1995-05-01   -0.008096
1995-06-01   -0.017395
dtype: float64
```

```
In [24]: predictions_ARIMA_diff_cumsum = predictions_ARIMA_diff.cumsum()
print (predictions_ARIMA_diff_cumsum.head())
```


```
Month
1995-02-01    0.011261
1995-03-01    0.027863
1995-04-01    0.049526
1995-05-01    0.041430
1995-06-01    0.024035
dtype: float64
```

```
In [25]: predictions_ARIMA_log = pd.Series(indexedDataset_logScale['Passengers'], index=indexedD
predictions_ARIMA_log = predictions_ARIMA_log.add(predictions_ARIMA_diff_cumsum, fill_va
predictions_ARIMA_log.head())
```

```
Out[25]: Month
1995-01-01    4.718499
1995-02-01    4.781946
1995-03-01    4.910665
1995-04-01    4.909338
1995-05-01    4.837220
dtype: float64
```

```
In [26]: predictions_ARIMA = np.exp(predictions_ARIMA_log)
plt.plot(indexedDataset)
plt.plot(predictions_ARIMA)
plt.title('RMSE: %.4f'% np.sqrt(sum((predictions_ARIMA-indexedDataset["Passengers"])**2
```

Out[26]: Text(0.5, 1.0, 'RMSE: 221.4940')



In []: