

```
In [1]: import pandas as pd
import numpy as np
```

```
In [5]: salarydata_train=pd.read_csv("C:\\Users\\Admin\\Downloads\\naive bayes\\SalaryData_Train.csv")
salarydata_train.head()
```

Out[5]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex	capitalgain	capitalloss	hoursperweek
0	39	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13
2	38	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40
3	53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40
4	28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40

```
In [7]: salarydata_test = pd.read_csv("C:\\Users\\Admin\\Downloads\\naive bayes\\SalaryData_Test.csv")
salarydata_test.head()
```

Out[7]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex	capitalgain	capitalloss	hoursperweek	n
0	25	Private	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0	0	40	U
1	38	Private	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0	0	50	U
2	28	Local-gov	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0	0	40	U
3	44	Private	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688	0	40	U
4	34	Private	10th	6	Never-married	Other-service	Not-in-family	White	Male	0	0	30	U

```
In [8]: salarydata_train.shape
```

```
Out[8]: (30161, 14)
```

```
In [9]: salarydata_test.shape
```

```
Out[9]: (15060, 14)
```

```
In [6]: salarydata_train.isnull().sum()
```

```
Out[6]: age                0  
workclass                0  
education                0  
educationno              0  
maritalstatus            0  
occupation               0  
relationship             0  
race                     0  
sex                      0  
capitalgain              0  
capitalloss              0  
hoursperweek             0  
native                   0  
Salary                   0  
dtype: int64
```

In [11]: salarydata_train.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30161 entries, 0 to 30160
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   30161 non-null  int64
1   workclass             30161 non-null  object
2   education             30161 non-null  object
3   educationno           30161 non-null  int64
4   maritalstatus         30161 non-null  object
5   occupation            30161 non-null  object
6   relationship          30161 non-null  object
7   race                  30161 non-null  object
8   sex                   30161 non-null  object
9   capitalgain           30161 non-null  int64
10  capitalloss           30161 non-null  int64
11  hoursperweek          30161 non-null  int64
12  native                30161 non-null  object
13  Salary                30161 non-null  object
dtypes: int64(5), object(9)
memory usage: 3.2+ MB
```

In [12]: salarydata_train.describe()

Out[12]:

	age	educationno	capitalgain	capitalloss	hoursperweek
count	30161.000000	30161.000000	30161.000000	30161.000000	30161.000000
mean	38.438115	10.121316	1092.044064	88.302311	40.931269
std	13.134830	2.550037	7406.466611	404.121321	11.980182
min	17.000000	1.000000	0.000000	0.000000	1.000000
25%	28.000000	9.000000	0.000000	0.000000	40.000000
50%	37.000000	10.000000	0.000000	0.000000	40.000000
75%	47.000000	13.000000	0.000000	0.000000	45.000000
max	90.000000	16.000000	99999.000000	4356.000000	99.000000

In [13]: salarydata_test.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15060 entries, 0 to 15059
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    15060 non-null  int64
1   workclass              15060 non-null  object
2   education              15060 non-null  object
3   educationno            15060 non-null  int64
4   maritalstatus          15060 non-null  object
5   occupation             15060 non-null  object
6   relationship           15060 non-null  object
7   race                   15060 non-null  object
8   sex                    15060 non-null  object
9   capitalgain            15060 non-null  int64
10  capitalloss            15060 non-null  int64
11  hoursperweek           15060 non-null  int64
12  native                  15060 non-null  object
13  Salary                 15060 non-null  object
dtypes: int64(5), object(9)
memory usage: 1.6+ MB
```

In [15]: salarydata_test.describe()

Out[15]:

	age	educationno	capitalgain	capitalloss	hoursperweek
count	15060.000000	15060.000000	15060.000000	15060.000000	15060.000000
mean	38.768327	10.112749	1120.301594	89.041899	40.951594
std	13.380676	2.558727	7703.181842	406.283245	12.062831
min	17.000000	1.000000	0.000000	0.000000	1.000000
25%	28.000000	9.000000	0.000000	0.000000	40.000000
50%	37.000000	10.000000	0.000000	0.000000	40.000000
75%	48.000000	13.000000	0.000000	0.000000	45.000000
max	90.000000	16.000000	99999.000000	3770.000000	99.000000

```
In [16]: salarydata_train.isin(['?']).sum(axis=0)
```

```
Out[16]: age                0
workclass                0
education                0
educationno              0
maritalstatus            0
occupation               0
relationship             0
race                     0
sex                      0
capitalgain              0
capitalloss              0
hoursperweek             0
native                   0
Salary                   0
dtype: int64
```

```
In [17]: salarydata_test.isin(['?']).sum(axis=0)
```

```
Out[17]: age                0
workclass                0
education                0
educationno              0
maritalstatus            0
occupation               0
relationship             0
race                     0
sex                      0
capitalgain              0
capitalloss              0
hoursperweek             0
native                   0
Salary                   0
dtype: int64
```

```
In [18]: print(salarydata_train[0:5])
```

	age	workclass	education	educationno	maritalstatus	\
0	39	State-gov	Bachelors	13	Never-married	
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	
2	38	Private	HS-grad	9	Divorced	
3	53	Private	11th	7	Married-civ-spouse	
4	28	Private	Bachelors	13	Married-civ-spouse	

	occupation	relationship	race	sex	capitalgain	\
0	Adm-clerical	Not-in-family	White	Male	2174	
1	Exec-managerial	Husband	White	Male	0	
2	Handlers-cleaners	Not-in-family	White	Male	0	
3	Handlers-cleaners	Husband	Black	Male	0	
4	Prof-specialty	Wife	Black	Female	0	

	capitalloss	hoursperweek	native	Salary
0	0	40	United-States	<=50K
1	0	13	United-States	<=50K
2	0	40	United-States	<=50K
3	0	40	United-States	<=50K
4	0	40	Cuba	<=50K

```
In [19]: categorical = [var for var in salarydata_train.columns if salarydata_train[var].dtype=='O']

print('There are {} categorical variables\n'.format(len(categorical)))

print('The categorical variables are :\n\n', categorical)
```

There are 9 categorical variables

The categorical variables are :

```
['workclass', 'education', 'maritalstatus', 'occupation', 'relationship', 'race', 'sex', 'native', 'Salary']
```

```
In [20]: salarydata_train[categorical].head()
```

```
Out[20]:
```

	workclass	education	maritalstatus	occupation	relationship	race	sex	native	Salary
0	State-gov	Bachelors	Never-married	Adm-clerical	Not-in-family	White	Male	United-States	<=50K
1	Self-emp-not-inc	Bachelors	Married-civ-spouse	Exec-managerial	Husband	White	Male	United-States	<=50K
2	Private	HS-grad	Divorced	Handlers-cleaners	Not-in-family	White	Male	United-States	<=50K
3	Private	11th	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	United-States	<=50K
4	Private	Bachelors	Married-civ-spouse	Prof-specialty	Wife	Black	Female	Cuba	<=50K

```
In [21]: salarydata_train[categorical].isnull().sum()
```

```
Out[21]: workclass      0
education    0
maritalstatus 0
occupation   0
relationship 0
race         0
sex          0
native       0
Salary       0
dtype: int64
```

In [22]: `for var in categorical:`

```
    print(salarydata_train[var].value_counts())
```

Nicaragua	33
Peru	30
Greece	29
France	27
Ecuador	27
Ireland	24
Hong	19
Cambodia	18
Trinidad&Tobago	18
Laos	17
Thailand	17
Yugoslavia	16
Outlying-US(Guam-USVI-etc)	14
Hungary	13
Honduras	12
Scotland	11

Name: native, dtype: int64

<=50K	22653
>50K	7508

Name: Salary, dtype: int64

In [23]: `for var in categorical:`

```
    print(salarydata_train[var].value_counts()/np.float(len(salarydata_train)))
```

```
Hong                0.000630
Cambodia            0.000597
Trinidad&Tobago     0.000597
Laos                0.000564
Thailand            0.000564
Yugoslavia          0.000530
Outlying-US(Guam-USVI-etc) 0.000464
Hungary             0.000431
Honduras            0.000398
Scotland            0.000365
```

Name: native, dtype: float64

```
<=50K    0.751069
```

```
>50K     0.248931
```

Name: Salary, dtype: float64

C:\Users\Admin\AppData\Local\Temp\ipykernel_432\217981199.py:3: DeprecationWarning: `np.float` is a deprecated alias for the builtin `float`. To silence this warning, use `float` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.float64` here. Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations> (<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>)

In [24]: `salarydata_train.workclass.unique()`

Out[24]: `array([' State-gov', ' Self-emp-not-inc', ' Private', ' Federal-gov',
 ' Local-gov', ' Self-emp-inc', ' Without-pay'], dtype=object)`

In [25]: `salarydata_train.workclass.value_counts()`

```
Out[25]: Private                22285
Self-emp-not-inc            2499
Local-gov                   2067
State-gov                   1279
Self-emp-inc                1074
Federal-gov                 943
Without-pay                 14
Name: workclass, dtype: int64
```

```
In [26]: salarydata_train.occupation.unique()
```

```
Out[26]: array([' Adm-clerical', ' Exec-managerial', ' Handlers-cleaners',  
                ' Prof-specialty', ' Other-service', ' Sales', ' Transport-moving',  
                ' Farming-fishing', ' Machine-op-inspct', ' Tech-support',  
                ' Craft-repair', ' Protective-serv', ' Armed-Forces',  
                ' Priv-house-serv'], dtype=object)
```

```
In [27]: salarydata_train.occupation.value_counts()
```

```
Out[27]: Prof-specialty      4038  
Craft-repair      4030  
Exec-managerial   3992  
Adm-clerical      3721  
Sales             3584  
Other-service     3212  
Machine-op-inspct 1965  
Transport-moving  1572  
Handlers-cleaners 1350  
Farming-fishing   989  
Tech-support      912  
Protective-serv   644  
Priv-house-serv   143  
Armed-Forces      9  
Name: occupation, dtype: int64
```

```
In [28]: salarydata_train.native.unique()
```

```
Out[28]: array([' United-States', ' Cuba', ' Jamaica', ' India', ' Mexico',  
                ' Puerto-Rico', ' Honduras', ' England', ' Canada', ' Germany',  
                ' Iran', ' Philippines', ' Poland', ' Columbia', ' Cambodia',  
                ' Thailand', ' Ecuador', ' Laos', ' Taiwan', ' Haiti', ' Portugal',  
                ' Dominican-Republic', ' El-Salvador', ' France', ' Guatemala',  
                ' Italy', ' China', ' South', ' Japan', ' Yugoslavia', ' Peru',  
                ' Outlying-US(Guam-USVI-etc)', ' Scotland', ' Trinidad&Tobago',  
                ' Greece', ' Nicaragua', ' Vietnam', ' Hong', ' Ireland',  
                ' Hungary'], dtype=object)
```

```
In [29]: salarydata_train.native.value_counts()
```

```
Out[29]: United-States      27504  
Mexico      610  
Philippines  188  
Germany     128  
Puerto-Rico 109  
Canada      107  
India        100  
El-Salvador  100  
Cuba         92  
England      86  
Jamaica      80  
South        71  
China        68  
Italy        68  
Dominican-Republic 67  
Vietnam      64  
Guatemala    63  
Japan        59  
Poland       56  
Columbia     56  
Iran         42  
Taiwan       42  
Haiti        42  
Portugal     34  
Nicaragua    33  
Peru         30  
Greece       29  
France       27  
Ecuador      27  
Ireland      24  
Hong         19  
Cambodia     18  
Trinidad&Tobago 18  
Laos         17  
Thailand     17  
Yugoslavia   16  
Outlying-US(Guam-USVI-etc) 14  
Hungary      13  
Honduras     12  
Scotland     11  
Name: native, dtype: int64
```

In [30]: `for var in categorical:`

```
    print(var, ' contains ', len(salarydata_train[var].unique()), ' labels')
```

```
workclass contains 7 labels
education contains 16 labels
maritalstatus contains 7 labels
occupation contains 14 labels
relationship contains 6 labels
race contains 5 labels
sex contains 2 labels
native contains 40 labels
Salary contains 2 labels
```

In [31]: `numerical = [var for var in salarydata_train.columns if salarydata_train[var].dtype != 'O']`

```
print('There are {} numerical variables\n'.format(len(numerical)))
```

```
print('The numerical variables are :', numerical)
```

There are 5 numerical variables

The numerical variables are : ['age', 'educationno', 'capitalgain', 'capitalloss', 'hoursperweek']

In [32]: `salarydata_train[numerical].head()`

Out[32]:

	age	educationno	capitalgain	capitalloss	hoursperweek
0	39	13	2174	0	40
1	50	13	0	0	13
2	38	9	0	0	40
3	53	7	0	0	40
4	28	13	0	0	40

```
In [33]: salarydata_train[numerical].isnull().sum()
```

```
Out[33]: age                0
educationno            0
capitalgain            0
capitalloss            0
hoursperweek          0
dtype: int64
```

```
In [34]: X = salarydata_train.drop(['Salary'], axis=1)
y = salarydata_train['Salary']
```

```
In [35]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 0)
```

```
In [36]: X_train.shape, X_test.shape
```

```
Out[36]: ((21112, 13), (9049, 13))
```

```
In [37]: X_train.dtypes
```

```
Out[37]: age                int64
workclass                object
education                object
educationno              int64
maritalstatus            object
occupation               object
relationship             object
race                    object
sex                    object
capitalgain              int64
capitalloss              int64
hoursperweek            int64
native                  object
dtype: object
```

```
In [38]: X_test.dtypes
```

```
Out[38]: age                int64
workclass                 object
education                 object
educationno               int64
maritalstatus             object
occupation                object
relationship              object
race                     object
sex                      object
capitalgain               int64
capitalloss               int64
hoursperweek              int64
native                   object
dtype: object
```

```
In [39]: categorical = [col for col in X_train.columns if X_train[col].dtypes == 'O']
```

```
categorical
```

```
Out[39]: ['workclass',
'education',
'maritalstatus',
'occupation',
'relationship',
'race',
'sex',
'native']
```

```
In [40]: numerical = [col for col in X_train.columns if X_train[col].dtypes != 'O']
```

```
numerical
```

```
Out[40]: ['age', 'educationno', 'capitalgain', 'capitalloss', 'hoursperweek']
```

```
In [41]: X_train[categorical].isnull().mean()
```

```
Out[41]: workclass      0.0  
education    0.0  
maritalstatus 0.0  
occupation   0.0  
relationship  0.0  
race         0.0  
sex          0.0  
native       0.0  
dtype: float64
```

```
In [42]: for col in categorical:  
         if X_train[col].isnull().mean()>0:  
             print(col, (X_train[col].isnull().mean()))
```

```
In [44]: X_train[categorical].isnull().sum()
```

```
Out[44]: workclass      0  
education    0  
maritalstatus 0  
occupation   0  
relationship  0  
race         0  
sex          0  
native       0  
dtype: int64
```

```
In [45]: X_test[categorical].isnull().sum()
```

```
Out[45]: workclass      0  
education    0  
maritalstatus 0  
occupation   0  
relationship  0  
race         0  
sex          0  
native       0  
dtype: int64
```

```
In [46]: X_train.isnull().sum()
```

```
Out[46]: age                0
workclass            0
education            0
educationno         0
maritalstatus       0
occupation          0
relationship        0
race                0
sex                 0
capitalgain         0
capitalloss         0
hoursperweek        0
native              0
dtype: int64
```

```
In [47]: X_test.isnull().sum()
```

```
Out[47]: age                0
workclass            0
education            0
educationno         0
maritalstatus       0
occupation          0
relationship        0
race                0
sex                 0
capitalgain         0
capitalloss         0
hoursperweek        0
native              0
dtype: int64
```



```
In [48]: categorical
```

```
Out[48]: ['workclass',  
          'education',  
          'maritalstatus',  
          'occupation',  
          'relationship',  
          'race',  
          'sex',  
          'native']
```

```
In [49]: X_train[categorical].head()
```

```
Out[49]:
```

	workclass	education	maritalstatus	occupation	relationship	race	sex	native
8166	Local-gov	Some-college	Married-civ-spouse	Protective-serv	Husband	White	Male	United-States
7138	Private	Some-college	Never-married	Other-service	Own-child	White	Male	United-States
437	Private	HS-grad	Never-married	Transport-moving	Not-in-family	White	Male	United-States
5436	Private	HS-grad	Married-civ-spouse	Craft-repair	Husband	White	Male	United-States
6541	Self-emp-not-inc	HS-grad	Married-civ-spouse	Tech-support	Husband	White	Male	United-States

```
In [50]: !pip install category_encoders
```

```
Requirement already satisfied: category_encoders in c:\users\admin\anaconda3\lib\site-packages (2.4.0)  
Requirement already satisfied: statsmodels>=0.9.0 in c:\users\admin\anaconda3\lib\site-packages (from category_encoders) (0.12.2)  
Requirement already satisfied: numpy>=1.14.0 in c:\users\admin\anaconda3\lib\site-packages (from category_encoders) (1.20.3)  
Requirement already satisfied: scikit-learn>=0.20.0 in c:\users\admin\anaconda3\lib\site-packages (from category_encoders) (0.24.2)  
Requirement already satisfied: pandas>=0.21.1 in c:\users\admin\anaconda3\lib\site-packages (from category_encoders) (1.3.4)  
Requirement already satisfied: patsy>=0.5.1 in c:\users\admin\anaconda3\lib\site-packages (from category_encoders) (0.5.2)  
Requirement already satisfied: scipy>=1.0.0 in c:\users\admin\anaconda3\lib\site-packages (from category_encoders) (1.7.1)  
Requirement already satisfied: pytz>=2017.3 in c:\users\admin\anaconda3\lib\site-packages (from pandas>=0.21.1->category_encoders) (2021.3)  
Requirement already satisfied: python-dateutil>=2.7.3 in c:\users\admin\anaconda3\lib\site-packages (from pandas>=0.21.1->category_encoders) (2.8.2)  
Requirement already satisfied: six in c:\users\admin\anaconda3\lib\site-packages (from patsy>=0.5.1->category_encoders) (1.16.0)  
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\admin\anaconda3\lib\site-packages (from scikit-learn>=0.20.0->category_encoders) (2.2.0)  
Requirement already satisfied: joblib>=0.11 in c:\users\admin\anaconda3\lib\site-packages (from scikit-learn>=0.20.0->category_encoders) (1.1.0)
```

```
In [51]: import category_encoders as ce
```

```
In [52]: encoder = ce.OneHotEncoder(cols=['workclass', 'education', 'maritalstatus', 'occupation', 'relationship',  
                                         'race', 'sex', 'native'])  
  
X_train = encoder.fit_transform(X_train)  
  
X_test = encoder.transform(X_test)
```

In [53]: `X_train.head()`

Out[53]:

	age	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5	workclass_6	workclass_7	education_1	education_2	...	r
8166	54	1	0	0	0	0	0	0	1	0	...	
7138	21	0	1	0	0	0	0	0	1	0	...	
437	30	0	1	0	0	0	0	0	0	1	...	
5436	42	0	1	0	0	0	0	0	0	1	...	
6541	37	0	0	1	0	0	0	0	0	1	...	

5 rows × 102 columns



In [55]: `X_train.shape`

Out[55]: (21112, 102)

In [56]: `X_test.head()`

Out[56]:

	age	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5	workclass_6	workclass_7	education_1	education_2	...	
25338	21	0	1	0	0	0	0	0	0	1	...	
18840	21	0	1	0	0	0	0	0	0	0	...	
8391	56	0	1	0	0	0	0	0	0	0	...	
18258	43	1	0	0	0	0	0	0	1	0	...	
16669	53	0	0	0	1	0	0	0	0	0	...	

5 rows × 102 columns



In [57]: `X_test.shape`

Out[57]: (9049, 102)

```
In [58]: cols = X_train.columns
```

```
In [59]: from sklearn.preprocessing import RobustScaler

scaler = RobustScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)
```

```
In [60]: X_train = pd.DataFrame(X_train, columns=[cols])
```

```
In [61]: X_test = pd.DataFrame(X_test, columns=[cols])
```

```
In [62]: X_train.head()
```

Out[62]:

	age	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5	workclass_6	workclass_7	education_1	education_2	...
0	0.894737	1.0	-1.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	...
1	-0.842105	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	...
2	-0.368421	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...
3	0.263158	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...
4	0.000000	0.0	-1.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	...

5 rows × 102 columns



```
In [63]: from sklearn.naive_bayes import GaussianNB

gnb = GaussianNB()

gnb.fit(X_train, y_train)
```

Out[63]: GaussianNB()

```
In [64]: y_pred = gnb.predict(X_test)
```

```
y_pred
```

```
Out[64]: array([' <=50K', ' <=50K', ' <=50K', ..., ' <=50K', ' <=50K', ' >50K'],  
          dtype='<U6')
```

```
In [65]: from sklearn.metrics import accuracy_score
```

```
print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test, y_pred)))
```

```
Model accuracy score: 0.7995
```

```
In [66]: y_pred_train = gnb.predict(X_train)
```

```
y_pred_train
```

```
Out[66]: array([' >50K', ' <=50K', ' <=50K', ..., ' <=50K', ' >50K', ' <=50K'],  
          dtype='<U6')
```

```
In [67]: print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train, y_pred_train)))
```

```
Training-set accuracy score: 0.8023
```

```
In [68]: print('Training set score: {:.4f}'.format(gnb.score(X_train, y_train)))
```

```
print('Test set score: {:.4f}'.format(gnb.score(X_test, y_test)))
```

```
Training set score: 0.8023
```

```
Test set score: 0.7995
```

```
In [69]: y_test.value_counts()
```

```
Out[69]: <=50K    6798  
>50K      2251  
Name: Salary, dtype: int64
```

```
In [70]: null_accuracy = (7407/(7407+2362))  
  
print('Null accuracy score: {0:0.4f}'.format(null_accuracy))  
  
Null accuracy score: 0.7582
```

```
In [71]: from sklearn.metrics import confusion_matrix  
  
cm = confusion_matrix(y_test, y_pred)  
  
print('Confusion matrix\n\n', cm)  
  
print('\nTrue Positives(TP) = ', cm[0,0])  
  
print('\nTrue Negatives(TN) = ', cm[1,1])  
  
print('\nFalse Positives(FP) = ', cm[0,1])  
  
print('\nFalse Negatives(FN) = ', cm[1,0])
```

Confusion matrix

```
[[5422 1376]  
 [ 438 1813]]
```

True Positives(TP) = 5422

True Negatives(TN) = 1813


False Positives(FP) = 1376

False Negatives(FN) = 438

```
In [73]: import seaborn as sns
cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'],
                          index=['Predict Positive:1', 'Predict Negative:0'])

sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
```

Out[73]: <AxesSubplot:>



```
In [74]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
<=50K	0.93	0.80	0.86	6798
>50K	0.57	0.81	0.67	2251
accuracy			0.80	9049
macro avg	0.75	0.80	0.76	9049
weighted avg	0.84	0.80	0.81	9049

```
In [75]: TP = cm[0,0]
TN = cm[1,1]
FP = cm[0,1]
FN = cm[1,0]
```

```
In [76]: classification_accuracy = (TP + TN) / float(TP + TN + FP + FN)
print('Classification accuracy : {0:0.4f}'.format(classification_accuracy))
Classification accuracy : 0.7995
```

```
In [77]: classification_error = (FP + FN) / float(TP + TN + FP + FN)
print('Classification error : {0:0.4f}'.format(classification_error))
Classification error : 0.2005
```

```
In [78]: precision = TP / float(TP + FP)
print('Precision : {0:0.4f}'.format(precision))
Precision : 0.7976
```



```
In [79]: recall = TP / float(TP + FN)

print('Recall or Sensitivity : {0:0.4f}'.format(recall))
```

Recall or Sensitivity : 0.9253

```
In [80]: true_positive_rate = TP / float(TP + FN)

print('True Positive Rate : {0:0.4f}'.format(true_positive_rate))
```

True Positive Rate : 0.9253

```
In [81]: false_positive_rate = FP / float(FP + TN)

print('False Positive Rate : {0:0.4f}'.format(false_positive_rate))
```

False Positive Rate : 0.4315

```
In [82]: specificity = TN / (TN + FP)

print('Specificity : {0:0.4f}'.format(specificity))
```

Specificity : 0.5685

```
In [83]: y_pred_prob = gnb.predict_proba(X_test)[0:10]

y_pred_prob
```

```
Out[83]: array([[9.99955511e-01, 4.44887598e-05],
 [9.95935549e-01, 4.06445120e-03],
 [8.63901480e-01, 1.36098520e-01],
 [9.99999906e-01, 9.37239455e-08],
 [8.80888343e-02, 9.11911166e-01],
 [9.99562896e-01, 4.37103927e-04],
 [5.34482750e-06, 9.99994655e-01],
 [6.28497161e-01, 3.71502839e-01],
 [5.46536963e-04, 9.99453463e-01],
 [9.9999570e-01, 4.30495598e-07]])
```

```
In [84]: y_pred_prob_df = pd.DataFrame(data=y_pred_prob, columns=['Prob of - <=50K', 'Prob of - >50K'])

y_pred_prob_df
```

Out[84]:

	Prob of - <=50K	Prob of - >50K
0	0.999956	4.448876e-05
1	0.995936	4.064451e-03
2	0.863901	1.360985e-01
3	1.000000	9.372395e-08
4	0.088089	9.119112e-01
5	0.999563	4.371039e-04
6	0.000005	9.999947e-01
7	0.628497	3.715028e-01
8	0.000547	9.994535e-01
9	1.000000	4.304956e-07

```
In [85]: gnb.predict_proba(X_test)[0:10, 1]
```

```
Out[85]: array([4.44887598e-05, 4.06445120e-03, 1.36098520e-01, 9.37239455e-08,
                9.11911166e-01, 4.37103927e-04, 9.99994655e-01, 3.71502839e-01,
                9.99453463e-01, 4.30495598e-07])
```

```
In [86]: y_pred1 = gnb.predict_proba(X_test)[: , 1]
```

```
In [89]: import matplotlib.pyplot as plt
plt.rcParams['font.size'] = 12


plt.hist(y_pred1, bins = 10)

plt.title('Histogram of predicted probabilities of salaries >50K')

plt.xlim(0,1)

plt.xlabel('Predicted probabilities of salaries >50K')
plt.ylabel('Frequency')
```

Out[89]: Text(0, 0.5, 'Frequency')



In [90]:

```
from sklearn.metrics import roc_curve

fpr, tpr, thresholds = roc_curve(y_test, y_pred1, pos_label = '>50K')

plt.figure(figsize=(6,4))

plt.plot(fpr, tpr, linewidth=2)

plt.plot([0,1], [0,1], 'k--' )

plt.rcParams['font.size'] = 12

plt.title('ROC curve for Gaussian Naive Bayes Classifier for Predicting Salaries')


plt.xlabel('False Positive Rate (1 - Specificity)')

plt.ylabel('True Positive Rate (Sensitivity)')

plt.show()
```

C:\Users\Admin\anaconda3\lib\site-packages\sklearn\metrics_ranking.py:949: UndefinedMetricWarning: No positive samples in y_true, true positive value should be meaningless
warnings.warn("No positive samples in y_true, "

ROC curve for Gaussian Naive Bayes Classifier for Predicting Salaries



```
In [91]: from sklearn.metrics import roc_auc_score  
  
ROC_AUC = roc_auc_score(y_test, y_pred1)  
  
print('ROC AUC : {:.4f}'.format(ROC_AUC))
```

ROC AUC : 0.8902

```
In [92]: from sklearn.model_selection import cross_val_score  
  
Cross_validated_ROC_AUC = cross_val_score(gnb, X_train, y_train, cv=5, scoring='roc_auc').mean()  
  
print('Cross validated ROC AUC : {:.4f}'.format(Cross_validated_ROC_AUC))
```

Cross validated ROC AUC : 0.8923

```
In [93]: from sklearn.model_selection import cross_val_score  
  
scores = cross_val_score(gnb, X_train, y_train, cv = 10, scoring='accuracy')  
  
print('Cross-validation scores:{}'.format(scores))
```

Cross-validation scores:[0.81676136 0.79829545 0.79014685 0.81288489 0.80388441 0.79062056
0.80767409 0.7925154 0.79630507 0.80909522]

```
In [94]: print('Average cross-validation score: {:.4f}'.format(scores.mean()))
```

Average cross-validation score: 0.8018

In []: